

# Optimization for Deep Learning

# Deep learning: recap


End-to-end joint learning of all layers:

- multiple assembleable blocks
- each block is piecewise-differentiable
- gradient-based **optimization**
- gradients computed by backpropagation



## Recap: training logistic regression

$$P(y(x) = y_i | w) = \frac{1}{1 + e^{-y_i w^T x_i}} = \sigma(y_i w^T x_i)$$



$$E(w) = - \sum_{i=1}^N \log P(y(x) = y_i | w) = \sum_{i=1}^N \log(1 + e^{-y_i w^T x_i})$$

# Softmax (sigmoid generalization)

Softmax (generalizes logistic):

$$\begin{array}{ccc} x & \xrightarrow[W^*]{\text{blue arrow}} & \begin{bmatrix} w_1^T x \\ w_2^T x \\ \dots \\ w_L^T x \end{bmatrix} \\ & & \downarrow \text{blue arrow } \textit{exp} \\ & & \begin{bmatrix} \exp(w_1^T x) \\ \exp(w_2^T x) \\ \dots \\ \exp(w_L^T x) \end{bmatrix} \\ & & \xrightarrow{\text{blue arrow } \textit{norm}} \end{array}$$

$P(y(x)=1)$   
//  
$$\begin{bmatrix} \frac{\exp(w_1^T x)}{\sum_i \exp(w_i^T x)} \\ \frac{\exp(w_2^T x)}{\sum_i \exp(w_i^T x)} \\ \dots \\ \frac{\exp(w_L^T x)}{\sum_i \exp(w_i^T x)} \end{bmatrix}$$

# Multinomial logistic loss

$$P(y(x) = i) = \frac{\exp w_i^T x}{\sum_j \exp w_j^T x}$$

Multinomial log loss (generalizes logistic loss):

$$\begin{aligned} E(w) &= - \sum_i \log P(y(x_i) = y_i) = \\ &= - \sum_i \left[ w_{y_i}^T x_i - \log \sum_j \exp w_j^T x_i \right] \end{aligned}$$

(Part of the) gradient over  $w_j$ :

$$\frac{dE}{dw_j} = - \sum_i x_i ([y_i == j] - P(y(x_i) = j))$$

# Sequential computation: *backpropagation*

$\frac{dz}{dx^3}, \frac{dz}{dw_4}$  can be computed

$$\frac{dz}{dw_3} = \frac{dx^3}{dw_3}^T \cdot \frac{dz}{dx^3}$$


$$\frac{dz}{dw_2} = \frac{dx^2}{dw_2}^T \cdot \frac{dz}{dx^2}$$

$$\frac{dz}{dw_1} = \frac{dx^1}{dw_1}^T \cdot \frac{dz}{dx^1}$$

$$\frac{dz}{dx^2} = \frac{dx^3}{dx^2}^T \cdot \frac{dz}{dx^3}$$

$$\frac{dz}{dx^1} = \frac{dx^2}{dx^1}^T \cdot \frac{dz}{dx^2}$$

$$\frac{dz}{dx^0} = \frac{dx^1}{dx^0}^T \cdot \frac{dz}{dx^1}$$



# Optimization for supervised ML

- $R(w)$  denotes regularization e.g.  $\|w\|^2$
- $l(x_i, y_i, w)$  denotes loss for  $i$ -th example, e.g.  $-\log P(y(x_i) = y_i | w)$
- The optimization objective is then:

$$E(w) = \frac{1}{N} \sum_{i=1}^N l(x_i, y_i, w) + \lambda R(w)$$

## Small scale setting: traditional optimization

$$E(w) = \frac{1}{N} \sum_{i=1}^N l(x_i, y_i, w) + \lambda R(w)$$

- Data are few, we can look through it at each optimization iteration
- Use adapted versions of standard optimization methods (gradient descent, quasi-Newton, quadratic programming,...)



# Large-scale learning

$$E(w) = \frac{1}{N} \sum_{i=1}^N l(x_i, y_i, w) + \lambda R(w)$$

$$\frac{dE}{dw} = \frac{1}{N} \sum_{i=1}^N \frac{dl(x_i, y_i, w)}{dw} + \lambda \frac{dR}{dw}$$

- Evaluating gradient is very expensive
- It will only be good for one (small) step

Stochastic gradient descent (SGD) idea:

- Evaluate a coarse approximation to grad
- Make “quick” steps

# Stochastic gradient descent (SGD)

Gradient:

$$\frac{dE}{dw} = \frac{1}{N} \sum_{i=1}^N \frac{dl(x_i, y_i, w)}{dw} + \lambda \frac{dR}{dw}$$

**Stochastic gradient:**

$$\frac{dE^i}{dw} = \frac{dl(x_i, y_i, w)}{dw} + \lambda \frac{dR}{dw}$$

Stochastic gradient is an unbiased estimate of the gradient:

$$\frac{dE}{dw} = \frac{1}{N} \sum_{i=1}^N \frac{dE^i}{dw}$$

# Stochastic gradient descent (SGD)

SGD:

$$\begin{aligned}v[t] &= -\alpha[t] \nabla(E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

where  $\nabla(E, w[t]) = \frac{dE^{i(t)}}{dw}(w[t])$

- $i(t)$  usually follow random permutations of training data
- One sweep over training data is called an **epoch**

# Stochastic gradient descent (SGD)

SGD:

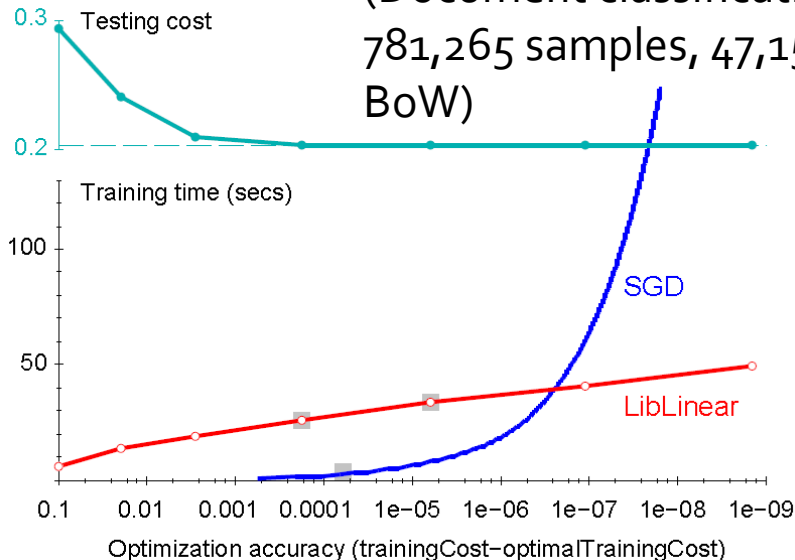
$$\begin{aligned}v[t] &= -\alpha[t] \nabla(E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

- One sweep over training data is called an **epoch**
- Popular choices for schedule  $\alpha[t]$ :
  - constant, e.g.  $\alpha[t] = 0.0001$
  - piecewise constant, e.g.  $\alpha[t]$  is decreased tenfold every  $N$  epochs
  - harmonic, e.g.  $\alpha[t] = 0.001 / ([t/N] + 10)$

# The efficiency of SGD ("shallow" learning)

[L.Bottou]

(Document classification :  
781,265 samples, 47,152-dim  
BoW)



# Batch SGD

Gradient:

$$\frac{dE}{dw} = \frac{1}{N} \sum_{i=1}^N \frac{dl(x_i, y_i, w)}{dw} + \lambda \frac{dR}{dw}$$

Batch (aka mini-batch):

$$\{b_1, b_2, \dots, b_{N_b}\} \subset 1 \dots N$$

Batch stochastic gradient:


$$\frac{dE}{dw} = \frac{1}{N_b} \sum_{i=1}^{N_b} \frac{dl(x_{b(i)}, y_{b(i)}, w)}{dw} + \lambda \frac{dR}{dw}$$

# Why do batching?

$$\frac{dE}{dw} = \frac{1}{N_b} \sum_{i=1}^{N_b} \frac{dl(x_{b(i)}, y_{b(i)}, w)}{dw} + \lambda \frac{dR}{dw}$$

- “Less stochastic” approximation, more stable convergence (questionable)
- **Main reason:** all modern architectures have parallelism, hence computing mini-batch grad is often as cheap as a single stochastic grad

# SGD inherits gradient descent problems





- Gradient descent is very poor “in ravines”
- SGD is no better



# Better optimization methods

- Second order methods (Newton, Quasi-Newton)
- Krylov subspace methods, in particular *conjugate gradients*



# Improving SGD using momentum

- Conjugate gradients use a combination of the current gradient and previous direction for the next step
- Similar idea for SGD (*momentum*):

$$\begin{aligned}v[t] &= -\alpha[t] \nabla(E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$



$$\begin{aligned}v[t] &= \mu v[t-1] - \alpha[t] \nabla(E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

Typical  $\mu = 0.9$

# Exponentially decaying running average


$$\begin{aligned}v[t] &= \mu v[t-1] - \alpha[t] \nabla (E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

$$\begin{aligned}v[t] &= \mu v[t-1] - \alpha[t] \nabla (E, w[t]) = \\&= \mu^2 v[t-2] - \mu \alpha[t-1] \nabla (E, w[t-1]) \\&\quad - \alpha[t] \nabla (f, w[t]) = \\&= \mu^3 v[t-3] - \mu^2 \alpha[t-2] \nabla (E, w[t-2]) \\&\quad - \mu \alpha[t-1] \nabla (E, w[t-1]) - \alpha[t] \nabla (E, w[t]) = \\&= \mu^{k+1} v[t-k-1] + \sum_{i=0}^k \mu^i \alpha[t-i] \nabla (E, w[t-i])\end{aligned}$$


# Momentum: why it works

$$v[t] \approx \sum_{i=0}^k \mu^i \alpha[t-i] \nabla(E, w[t-i])$$

- Smoothes out noise in SGD (~bigger batches)
- **Smoothes out oscillations inherent to gradient descent**
- Escapes local minima



# The effect of the momentum



Step-size  $\alpha = 0.0051$




Momentum  $\beta = 0.90$



We often think of Momentum as a means of speeding up the iterations, leading to faster convergence. However, it can also exhibit other interesting behavior. It allows a larger step-size to be used, and creates its own oscillations. What happens if we use a larger step-size?

[Goh, Distill 2017]

# Phase space along a single eigenvector



Ripples



$R$ 's eigenvalues are complex, and the iterates display low frequency ripples. Surprisingly, the convergence rate  $2\sqrt{\beta}$  is independent of  $\alpha$  and  $\lambda_i$ .

Monotonic Decrease




$R$ 's eigenvalues are both real, are positive, and have norm less than one. The behavior here resembles gradient descent.

1-Step Convergence




When  $\alpha = 1/\lambda_i$ , and  $\beta = 0$ , we converge in one step. This is a very special point, and kills the error in the eigenspace completely.

Monotonic Oscillations



When  $\alpha > 1/\lambda_i$ , the iterates flip between  $+$  and  $-$  at each iteration. These are often referred to as 'oscillations' in gradient descent.


Divergence



When  $\max\{|\sigma_1|, |\sigma_2|\} > 1$ , the iterates diverge.

[Goh, Distill 2017]

# Momentum: multiple eigenvalues



Optimal rate & momentum:

$$\alpha = \left( \frac{2}{\sqrt{\lambda_1} + \sqrt{\lambda_n}} \right)^2 \quad \beta = \left( \frac{\sqrt{\lambda_n} - \sqrt{\lambda_1}}{\sqrt{\lambda_n} + \sqrt{\lambda_1}} \right)^2$$

Optimal speed:

$$\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1}$$

Convergence rate,  
**Momentum**

$$\frac{\kappa - 1}{\kappa + 1}$$

Convergence rate,  
**Gradient Descent**

[Goh, Distill 2017]

# Momentum: multiple eigenvalues

Optimal rate &  
momentum:

$$\alpha = \left( \frac{2}{\sqrt{\lambda_1} + \sqrt{\lambda_n}} \right)^2 \quad \beta = \left( \frac{\sqrt{\lambda_n} - \sqrt{\lambda_1}}{\sqrt{\lambda_n} + \sqrt{\lambda_1}} \right)^2$$

In real network we do not know eigenvalues, so:

- we set the momentum high (e.g. 0.9)
- then we tune the learning rate



# Nesterov accelerated gradient

$$\begin{aligned}v[t] &= \mu v[t-1] - \alpha[t] \nabla(E, w[t]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

Before we even compute the gradient, we have a good approximation where we will end up:  $w[t+1] \approx w[t] + \mu v[t-1]$

Let us use this knowledge:

$$\begin{aligned}v[t] &= \mu v[t-1] - \alpha[t] \nabla(E, w[t] + \mu v[t-1]) \\w[t+1] &= w[t] + v[t]\end{aligned}$$

(Computing the gradient at a more relevant spot)

## Second-order methods

- Exponential smoothing helps, but still not optimal if large anisotropy exists
- Classic (Newton) solution: estimate the Hessian and make the update
$$v[t+1] = -H[t]^{-1} \nabla(E, w[t])$$
(the lower the curvature the faster we go)
- Quasi-Newton methods: estimate some approximation to Hessian based on observed gradients
- Quasi-Newton can be used in batch mode, but same batch should be used over several iterations (why?)

## Adagrad method [Duchi et al. 2011]

Adagrad idea: scale updates along different dimensions according to accumulated gradient magnitude

$$g[t] = g[t-1] + \nabla(E, w[t]) \odot \nabla(E, w[t])$$


$$w[t+1] = w[t] - \frac{\alpha}{\sqrt{g[t] + \epsilon}} \odot \nabla(E, w[t])$$

Note: step lengths automatically decrease (perhaps too quickly).

## Adagrad method [Duchi et al. 2011]

$$g[t] = g[t-1] + \nabla(E, w[t]) \odot \nabla(E, w[t])$$

$$w[t+1] = w[t] - \frac{\alpha}{\sqrt{g[t] + \epsilon}} \odot \nabla(E, w[t])$$



Adagrad in this case: find out that “vertical” derivatives are bigger, then make “vertical” steps smaller than “horizontal”

## RMSPROP method [Hinton 2012]

Same as Adagrad, but replace accumulation of squared gradient with running averaging:

$$g[t] = \mu g[t-1] + (1-\mu) \nabla(E, w[t]) \odot \nabla(E, w[t])$$

$$w[t+1] = w[t] - \frac{\alpha[t]}{\sqrt{g[t]} + \epsilon} \odot \nabla(E, w[t])$$

# Comparison: logistic regression




Image credit: Alec Redford

# Further comparison




Image credit: Alec Redford

# Further comparison: escaping from a saddle




Image credit: Alec Redford



# ADAM method [Kingma & Ba 2015]

ADAM = "ADaptive Moment Estimation"

$$v[t] = \beta v[t-1] + (1 - \beta) \nabla(E, w[t])$$


$$g[t] = \mu g[t-1] + (1 - \mu) \nabla(E, w[t]) \odot \nabla(E, w[t])$$

$$w[t+1] = w[t] - \alpha \frac{1}{\sqrt{g[t] + \epsilon}} \odot v[t]$$



*(Note: In the original image, blue lines indicate that the term  $\frac{1}{\sqrt{g[t] + \epsilon}}$  is multiplied by  $1 - \mu^t$  and the term  $v[t]$  is multiplied by  $1 - \beta^t$ .)*

Recommended values:  $\beta = 0.9$ ,  $\mu = 0.999$ ,  $\alpha = 0.001$ ,  $\epsilon = 10^{-8}$

# ADAM method [Kingma & Ba 2015]




# ADAM method [Kingma & Ba 2015]



# Recap: optimization methods for DL


- Stochastic optimization is used always
- Optimization methods are not trying to estimate full Hessian (ignoring interaction between variables)



Toy example:  $z = w_4 w_3 w_2 w_1 x$

$$\frac{dz}{dw_2} = w_4 w_3 w_1 x$$

# Problems with DL optimization




Toy example:  $z = w_4 w_3 w_2 w_1 x$

$$\frac{dz}{dw_2} = w_4 w_3 w_1 x$$

$$\frac{dz}{dw_3} = w_4 w_2 w_1 x$$


- $w = (1, 1, 1, 1)$  and  $w = (1, 0.01, 100, 1)$  define the same function ("gauge freedom"), but very different derivatives
- In the first case, derivatives and values are of order 1.
- In the second case, derivatives and values are wildly different

# Gauge freedom in ReLu Networks



$$\alpha > 0$$

$$1/\alpha \text{ ReLu}(\alpha x) = \text{ReLu}(x)$$



**Thus:** we can easily construct ReLu networks with **different** weights implementing the **same** function

# Normalizing in the toy example




Toy example: 
$$z_i = w_4 w_3 \frac{w_2 w_1 x_i}{\frac{1}{N} \sum_{j=1}^N w_2 w_1 x_j}$$

$$\left[ \frac{dz}{dw_3} \right]^i = w_4 \frac{w_2 w_1 x_i}{\frac{1}{N} \sum_{j=1}^N w_2 w_1 x_j}$$

- Now, increasing  $w_2$  or  $w_1$  100x times will not change the partial derivative w.r.t.  $w_3$  !
- The learning will become more stable

# Batch normalization

[Szegedy and Ioffe 2015]



- Makes the training process invariant to some re-parameterizations
- Eliminates the bulk of cross-layer correlation between derivatives (off-diagonal Hessian vals)
- Use mini-batch statistics at training time to ensure that neuron activations are distributed “nicely” and the learning proceeds



# Batch normalization layer

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1...m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

covariant to reparameterization →  $\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$  // mini-batch mean

→  $\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$  // mini-batch variance

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

learnable by SGD

[Szegedy and Ioffe 2015]

# Batch normalization layer

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

- At training time mean and variance are estimated per batch
- At test time, usually (running) averages over the dataset are used
- At test time, batch norm can be “merged in”
- For small batches, this is a big test $\leftrightarrow$ train mismatch ☹

Solutions to train-test mismatch:

- Keep training time behavior
- Switch to test behavior and fine-tune

# Alternatives to BatchNorm

- Layer Norm [Ba et al. NIPS'16], Instance Norm [Ulyanov et al. Arxiv16], Group renorm [Wu and He, ECCV18] – normalize over statistics of certain specific groups of variables **within** the same sample
- Batch Renorm [Ioffe NIPS'17]: gradually switch between train and test time behavior during training
- Weight norm [Salimans and Kingma NIPS'16]: decouple direction and magnitude of weight matrices

# Initialization schemes

- **Basic idea 1:** units should be initialized to have comparable total input weights
- **Basic idea 2:** use layers which keep magnitude (otherwise both forwardprop and backprop will suffer from explosion/attenuation to zero; normalization layers solve this issue)
- E.g. [Glorot&Bengio 2010] aka “Xavier-initialization”:

$$W \sim U\left[-\frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}, \frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}\right]$$

- E.g. [He et al, Arxiv15] for ReLu networks:

$$W \sim \mathcal{N}(0, \sqrt{2/n_i})$$

# Recap

- Batch SGD optimization is used in large-scale setting
- Advanced SGD methods use running averages to smooth and rescale SGD steps
- Normalization layers are important and used in most modern deep architectures

# Bibliography

Léon Bottou, Olivier Bousquet:

The Tradeoffs of Large Scale Learning. NIPS 2007: 161-168

Nesterov, Yurii. "A method of solving a convex programming problem with convergence rate  $O(1/k^2)$ ." Soviet Mathematics Doklady. Vol. 27. No. 2. 1983.

G. Goh, Why momentum really works? DISTILL 2017

<https://distill.pub/2017/momentum/>

John C. Duchi, Elad Hazan, Yoram Singer:

Adaptive Subgradient Methods for Online Learning and Stochastic Optimization. Journal of Machine Learning Research 12: 2121-2159 (2011)

Matthew D. Zeiler:

ADADELTA: An Adaptive Learning Rate Method. CoRR abs/1212.5701

Kingma, Diederik, and Jimmy Ba. "Adam: A method for stochastic optimization." ICLR 2015

# Bibliography

Sergey Ioffe, Christian Szegedy:

Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift.

ICML2015: 448-456

Sergey Ioffe, Batch Renormalization. NIPS 2017

Dmitry Ulyanov, Vadim Lebedev, Andrea Vedaldi, Victor S. Lempitsky:

Texture Networks: Feed-forward Synthesis of Textures and Stylized Images. ICML 2016: 1349-1357

Lei Jimmy Ba, Jamie Ryan Kiros, Geoffrey E. Hinton:

Layer Normalization. CoRR abs/1607.06450 (2016)

Yuxin Wu, Kaiming He: Group Normalization. ECCV (13) 2018: 3-19

Tim Salimans, Diederik P. Kingma:

Weight Normalization: A Simple Reparameterization to Accelerate Training of Deep Neural Networks.

NIPS 2016: 901

Xavier Glorot, Yoshua Bengio:

Understanding the difficulty of training deep feedforward neural networks. AISTATS 2010: 249-256

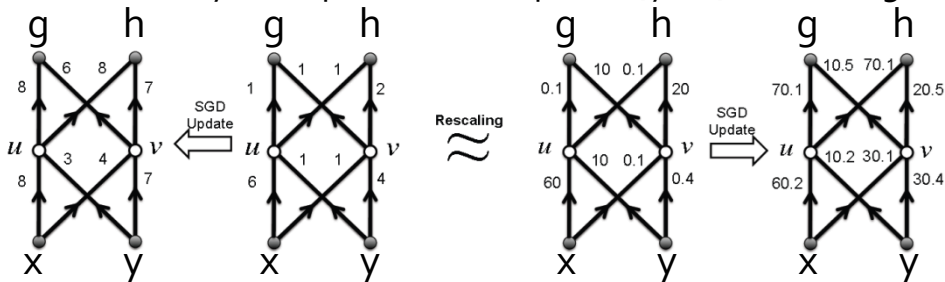
Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun:

Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification. ICCV

2015: 1026-1034

# Initialization schemes

One more toy example: 1 SGD step for  $(x, y = 1, 1)$  and  $L = g + h$



[Neyshabur, Salakhutdinov, Srebro, Path-SGD: Path-Normalized Optimization in Deep Neural Networks, NIPS2015]





# Units of measurements

- Let our coordinates be measured in meters. What is the unit of measurement for gradients?  
Assume unitless function...
- (Stochastic) gradient descent is inconsistent.
- Newton method is consistent.



## Adadelta method [Zeiler 2012]

$$g[t] = \mu g[t-1] + (1-\mu) \nabla(E, w[t]) \odot \nabla(E, w[t])$$

$$w[t+1] = w[t] - \frac{\sqrt{d[t]} + \epsilon}{\sqrt{g[t]} + \epsilon} \odot \nabla(E, w[t])$$

$$d[t+1] = \mu d[t] + (1-\mu) (w[t+1]-w[t]) \odot (w[t+1]-w[t])$$

- No step length parameter (good!)
- Correct units within the updates